

Artificial Neural Network Application for Reservoir Characterization Code ▼

Chukwuemeka Okoli (<https://www.linkedin.com/in/chukwuemeka-okoli-38686923/>)

April, 09, 2019

- Introduction
- The Goal
- Preliminaries
 - Data Import
 - Summaries
- Data Analysis
- Model Interpretability
- Conclusion

Introduction

As a Data Scientist in the Oil and Gas Industry, we are often faced with numerous challenges. For instance, with available Gas production data for an unconventional reservoir, can we establish a relationship between several independent variable in the production process and the dependent variable “Gas production per months” using Artificial Neural Network. In this R Markdown (<https://rmarkdown.rstudio.com>) project, an illustration of how we can apply Artificial Neural Network to Gas production is discussed. The dataset for this project is the “Q4_data.csv” file.

The Goal

Find a relationship between the dependent variable “Gas production per months” and several independent variables. Essentially, we want to determine the Gas Production per month for the given well data based on different factors.

Preliminaries

First, the required packages are installed using the `install.packages()` function.

Hide

```
# install knitr
install.packages("knitr")

# install the rmarkdown package
install.packages("rmarkdown")

# install ggplot for plotting
install.packages("ggplot2")

# install dplyr for data manipulation
install.packages("dplyr")

# install neuralnet for artificial neural network
install.packages("neuralnet")
install.packages("NeuralNetTools")
```

The installed package is then loaded using the `library()` function.

[Hide](#)

```
# load libraries
library(knitr)
library(rmarkdown)
library(ggplot2)
library(dplyr)
library(neuralnet)
library(NeuralNetTools)
library(xfun)
```

Data Import

We can call in the dataset into R using the code below. The data is stored in a `.csv` file.

[Hide](#)

```
set.seed(440)
data <- read.csv("Q4_data.csv", header=TRUE)
```

Summaries

The data object `field_data` is complex. It contains various information about properties of the well such as the perforation interval, fracture volume, proppant, casing depth, etc. Use `summary()` to get a quick summary of the data.

[Hide](#)

```
summary(data)
```

Perf_Int	Frac_vol	Proppant	N.frac	Tubing.depth	Casing.d	
epth						
Min. : 12.0	Min. : 7478	Min. : 43500	Min. : 1.000	Min. : 4952	Min. :	
7325						
1st Qu.: 381.8	1st Qu.: 29524	1st Qu.: 164625	1st Qu.: 1.000	1st Qu.: 7652	1st Qu.:	
8176						
Median : 474.0	Median : 36200	Median : 296946	Median : 2.000	Median : 8002	Median :	
8745						
Mean : 708.1	Mean : 56848	Mean : 662449	Mean : 2.309	Mean : 8020	Mean :	
9044						
3rd Qu.: 802.5	3rd Qu.: 56876	3rd Qu.:1135250	3rd Qu.: 3.000	3rd Qu.: 8407	3rd Qu.:	
9694						
Max. :4555.0	Max. :3637776	Max. :4437897	Max. :11.000	Max. :12100	Max. :1	
3144						
FTP	Choke.Size	SITHP	SG_gas	Well_Type	Latitude	
Min. : 0.0	Min. : 6.00	Min. : 300	Min. :0.5810	Min. :1.000	Min. :	
33						
1st Qu.: 450.0	1st Qu.: 14.00	1st Qu.:1615	1st Qu.:0.7000	1st Qu.:2.000	1st Qu.:	
33						
Median : 814.1	Median : 20.00	Median :2278	Median :0.7000	Median :3.000	Median :	
33						
Mean : 905.6	Mean : 28.18	Mean :2250	Mean :0.6922	Mean :2.504	Mean : 77	
37						
3rd Qu.:1265.0	3rd Qu.: 32.00	3rd Qu.:2913	3rd Qu.:0.7000	3rd Qu.:3.000	3rd Qu.:	
33						
Max. :2964.0	Max. :128.00	Max. :3715	Max. :1.0320	Max. :3.000	Max. :34361	
28						
Long.	Gas_prod	Acid				
Min. : -97.40	Min. : 95	Min. :1.000				
1st Qu.: -97.35	1st Qu.: 3840	1st Qu.:1.000				
Median : -97.31	Median : 6392	Median :1.000				
Mean : -97.00	Mean : 13281	Mean :1.193				
3rd Qu.: -97.26	3rd Qu.: 13155	3rd Qu.:1.000				
Max. : 33.04	Max. :1189790	Max. :2.000				

The `set.seed()` function is useful when running simulations to ensure all results, figures, etc. are reproducible. We can then check that no data point is missing. If we have a missing data point, we need to fix the dataset.

Hide

```
# Check that no data is missing
apply(data,2,function(x) sum(is.na(x)))
```

Perf_Int	Frac_vol	Proppant	N.frac	Tubing.depth	Casing.depth	FTP	Cho
ke.Size							
0	0	0	0	0	0	0	
0							
SITHP	SG_gas	Well_Type	Latitude	Long.	Gas_prod	Acid	
0	0	0	0	0	0	0	

We can see that there is no missing data. In the case where missing data exist, data cleaning will have to take place.

Data Analysis

To analyze the data, we first proceed by randomly splitting the data into a *train* and a *test* dataset. A linear regression model is fit to the train dataset, and tested on the test dataset.

Hide

```
# Train-test random splitting for linear model
index <- sample(1:nrow(data), round(0.75*nrow(data)))
train <- data[index,]
test <- data[-index,]

# Fitting linear model to the train dataset
lm.fit <- lm(Gas_prod~., data=train)
summary(lm.fit)
```

```
Call:
lm(formula = Gas_prod ~ ., data = train)

Residuals:
    Min       1Q   Median       3Q      Max
-47265  -10594   -1262    6029  1130490

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.242e+06  6.496e+06   0.653  0.51423
Perf_Int      7.347e+00  9.449e+00   0.778  0.43739
Frac_vol     -2.151e-03  1.828e-02  -0.118  0.90643
Proppant     -2.850e-03  8.633e-03  -0.330  0.74148
N.frac        3.468e+03  3.381e+03   1.026  0.30573
Tubing.depth  2.724e+00  7.987e+00   0.341  0.73332
Casing.depth -4.627e+00  7.114e+00  -0.650  0.51589
FTP           1.334e+01  7.556e+00   1.766  0.07839 .
Choke.Size   -1.336e+02  1.828e+02  -0.731  0.46519
SITHP        -1.010e+01  5.776e+00  -1.749  0.08126 .
SG_gas       2.025e+04  8.302e+04   0.244  0.80744
Well_Type    -1.712e+04  5.909e+03  -2.897  0.00402 **
Latitude     -1.619e+00  2.527e+00  -0.641  0.52226
Long.        4.283e+04  6.657e+04   0.643  0.52049
Acid         -7.482e+03  9.953e+03  -0.752  0.45275
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 64930 on 319 degrees of freedom
Multiple R-squared:  0.05826,    Adjusted R-squared:  0.01693
F-statistic:  1.41 on 14 and 319 DF,  p-value: 0.1465
```

Hide

```
# Predicted data from linear model fit
pr.lm <- predict(lm.fit, test)

# Test MSE
MSE.lm <- sum((pr.lm - test$Gas_prod)^2)/nrow(test)
```

Before fitting a neural network, some preparation needs to be done. Neural networks are not that easy to train and tune. As a first step, we need to address data preprocessing. It is good practice to normalize your data before training a neural network. This step is important because, depending on your dataset, avoiding normalization may lead to useless results or to a very difficult training process (most of the time, the algorithm will not converge before the number of maximum iterations allowed). You can choose different methods to scale the data (z-normalization, min-max scale, etc.). The data were scaled using the min-max method and scaled in the interval [0,1]. We therefore scale and split the data before moving on.

Hide

```
# Neural network fitting

# Scaling data for the Neural Network
maxs <- apply(data, 2, max)
mins <- apply(data, 2, min)
scaled <- as.data.frame(scale(data, center = mins, scale = maxs - mins))

# Train-test split
train_ <- scaled[index,]
test_ <- scaled[-index,]
```

Since there is no fixed rule as to how many layers and neurons to use, we are going to use two (2) hidden layers with 5 and 3 neurons. To fit the network, we use the following code:

Hide

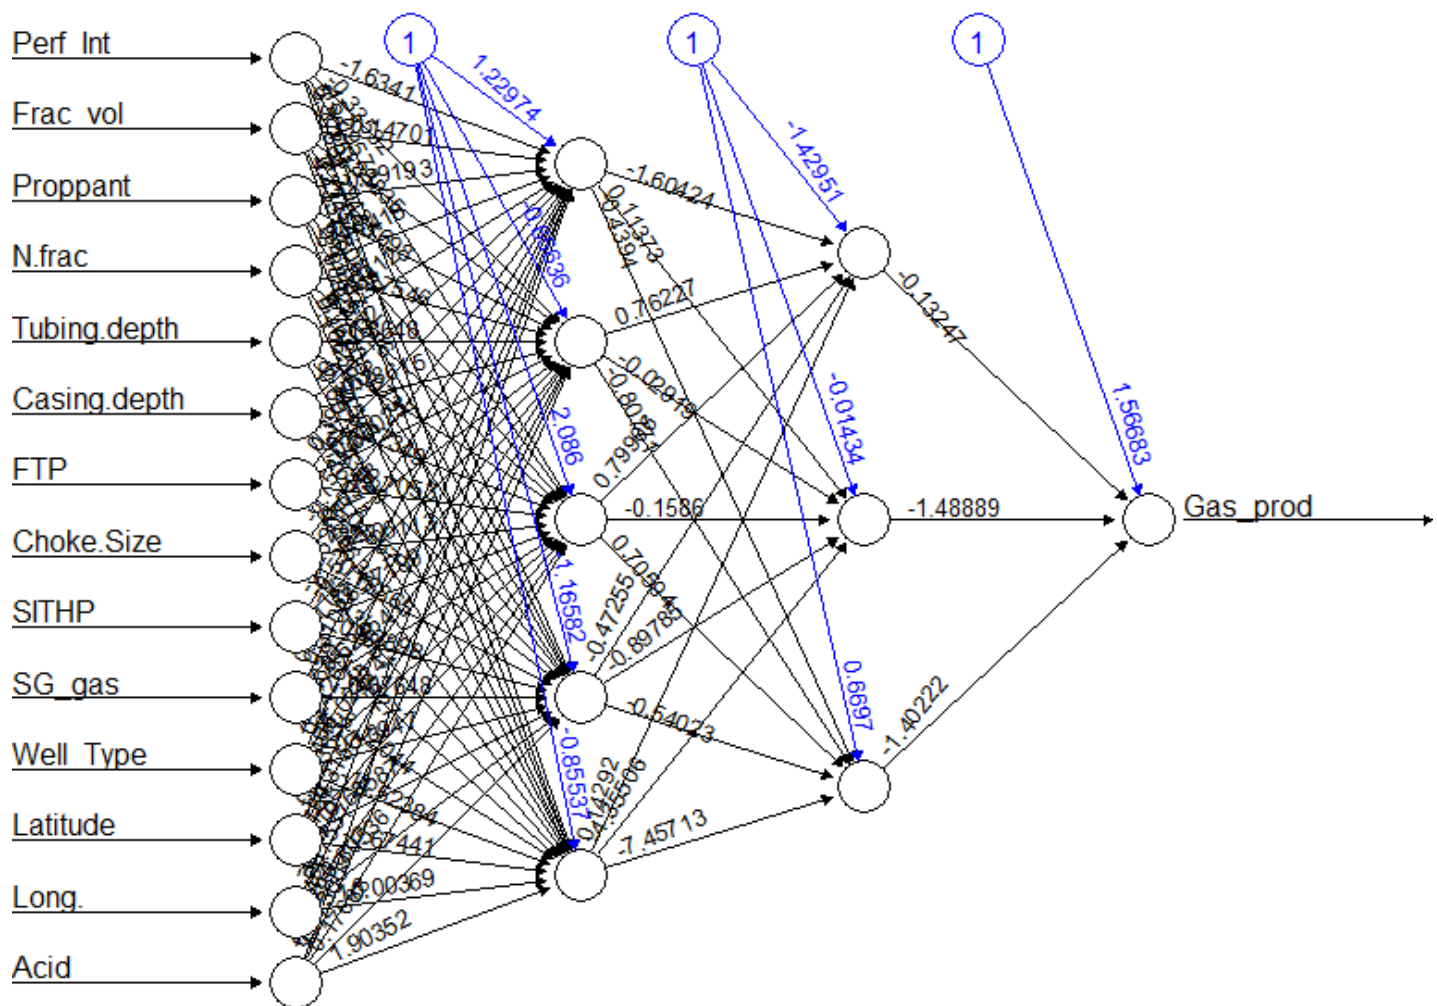
```
# Training the neural network
n <- names(train_)
f <- as.formula(paste("Gas_prod ~", paste(n[!n %in% "Gas_prod"], collapse = " + ")))
neural_net <- neuralnet(f, data = train_, hidden = c(5,3), linear.output = T)
```

Note that the hidden argument accepts a vector with the number of neurons for each hidden layer, while the argument linear.output is used to specify whether we want to do regression **linear.output = TRUE** or **classification linear.output = FALSE**

The neuralnet package provides a nice tool to plot the model. Use the following code to plot the neural network in R:

Hide

```
# Visual plot of the model
plot(neural_net)
```



Now we can try to predict the values for the test set and calculate the mean squared error (MSE). Remember that the net will output a normalized prediction, so we need to scale it back in order to make a meaningful comparison (or just a simple prediction). The mean squared error is one metric used to measure prediction accuracy. The MSE is calculated as:

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

[Hide](#)

```
# Prediction
pr.nn <- compute(neural_net, test_[,1:14])

# Results from Neural Networks are normalized (scaled)
# Descaling for comparison
pr.nn_ <- pr.nn$net.result*(max(data$Gas_prod)-min(data$Gas_prod))+min(data$Gas_prod)
test.r <- (test_$Gas_prod)*(max(data$Gas_prod)-min(data$Gas_prod))+min(data$Gas_prod)

# Calculating MSE
MSE.nn <- sum((test.r - pr.nn_)^2)/nrow(test_)
```

We then compare the two mean squares error obtained from the linear regression and the neural network fitting process using the code:

Hide

```
# Compare the two Mean Squared Errors (MSEs)
print(paste(MSE.lm, MSE.nn))
```

```
[1] "183730606.293678 132133926.405935"
```

This shows that the network is doing a better job at predicting Gas_prod than the linear model. The prediction from the neural network is **129413554.366771** which is better than the **229551018.949411** obtained from the linear model. We can perform a fast cross visualization in order to be more confident of the result. A visual approach to the performance of the network and the linear model is plotted below.

Hide

```
# Plot predictions
par(mfrow=c(1,2))

plot(test$Gas_prod, pr.nn_, col='red', main='Real vs predicted NN', pch=18, cex=0.7)
abline(0,1,lwd=2)
```


Hide

```
legend('bottomright', legend='NN', pch=18, col='red', bty='n')

plot(test$Gas_prod, pr.lm, col='blue', main='Real vs predicted lm', pch=18, cex=0.7)
```

Hide

```
abline(0,1,lwd=2)
legend('bottomright', legend='LM', pch=18, col='blue', bty='n', cex=.95)
```




By visually inspecting the plot, we can see that the predictions made by the neural network are more concentrated around the line than those made by the linear model. We can obtain a more useful visual comparison using the code below

[Hide](#)

```
# Compare predictions on the same plot
plot(test$Gas_prod, pr.nn_, col='red', main='Real vs predicted NN', pch=18, cex=0.7)
points(test$Gas_prod, pr.lm, col='blue', pch=18, cex=0.7)
```

[Hide](#)

```
abline(0,1,lwd=2)
legend('bottomright', legend=c('NN','LM'), pch=18, col=c('red','blue'))
```

The comparison of prediction on the same plot is shown above. In analyzing the model above, we allowed the network select the training, validation and testing dataset itself by randomly splitting the data into a train and a test set, then a linear regression model is fit and then tested on the test set.

If we choose to select the training, validation and testing datasets ourselves, we normalize our data and split into training and test data as before.

[Hide](#)

```
# Neural net fitting
set.seed(440)
mydata <- read.csv("Q4_data.csv", header=TRUE)

#Normalizing the dataset
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x)))
}

maxmindf <- as.data.frame(lapply(mydata, normalize))
```

[Hide](#)

```
# Display head of data
head(mydata)
```

	Perf_Int <int>	Frac_vol <dbl>	Proppant <int>	N.frac <int>	Tubing.depth <dbl>	Casing.depth <dbl>	FTP <dbl>	Choke.Size <dbl>	SI... <dbl>	
1	280	19454.76	408550	2	7754	8067	185.0	42	2815	
2	757	58477.00	407000	1	8843	9081	490.0	40	1915	
3	532	27262.00	129400	1	7963	8515	565.0	22	3015	
4	345	31622.74	321300	1	7832	8039	615.0	20	2350	
5	856	18596.76	390532	2	8602	9065	715.0	24	3250	
6	470	34785.71	190000	1	7470	7810	1235.2	10	2250	

6 rows | 1-10 of 15 columns

Hide

```
# Display tail of data
tail(mydata)
```

	Perf_Int <int>	Frac_vol <dbl>	Proppant <int>	N.frac <int>	Tubing.depth <dbl>	Casing.depth <dbl>	F... <dbl>	Choke.Size <dbl>	SI... <dbl>	
441	718	52929.00	701000	1	8851	8869	700	32	2200	
442	575	28945.00	465800	2	8115	8532	680	24	2450	
443	786	52379.00	352572	3	8426	9151	1365	20	2910	
444	740	23752.00	203500	2	8149	9200	439	18	2714	
445	702	24676.00	294765	2	8484	9245	1365	14	3070	
446	372	28928.57	115250	1	7370	7644	1500	9	2865	

6 rows | 1-10 of 15 columns

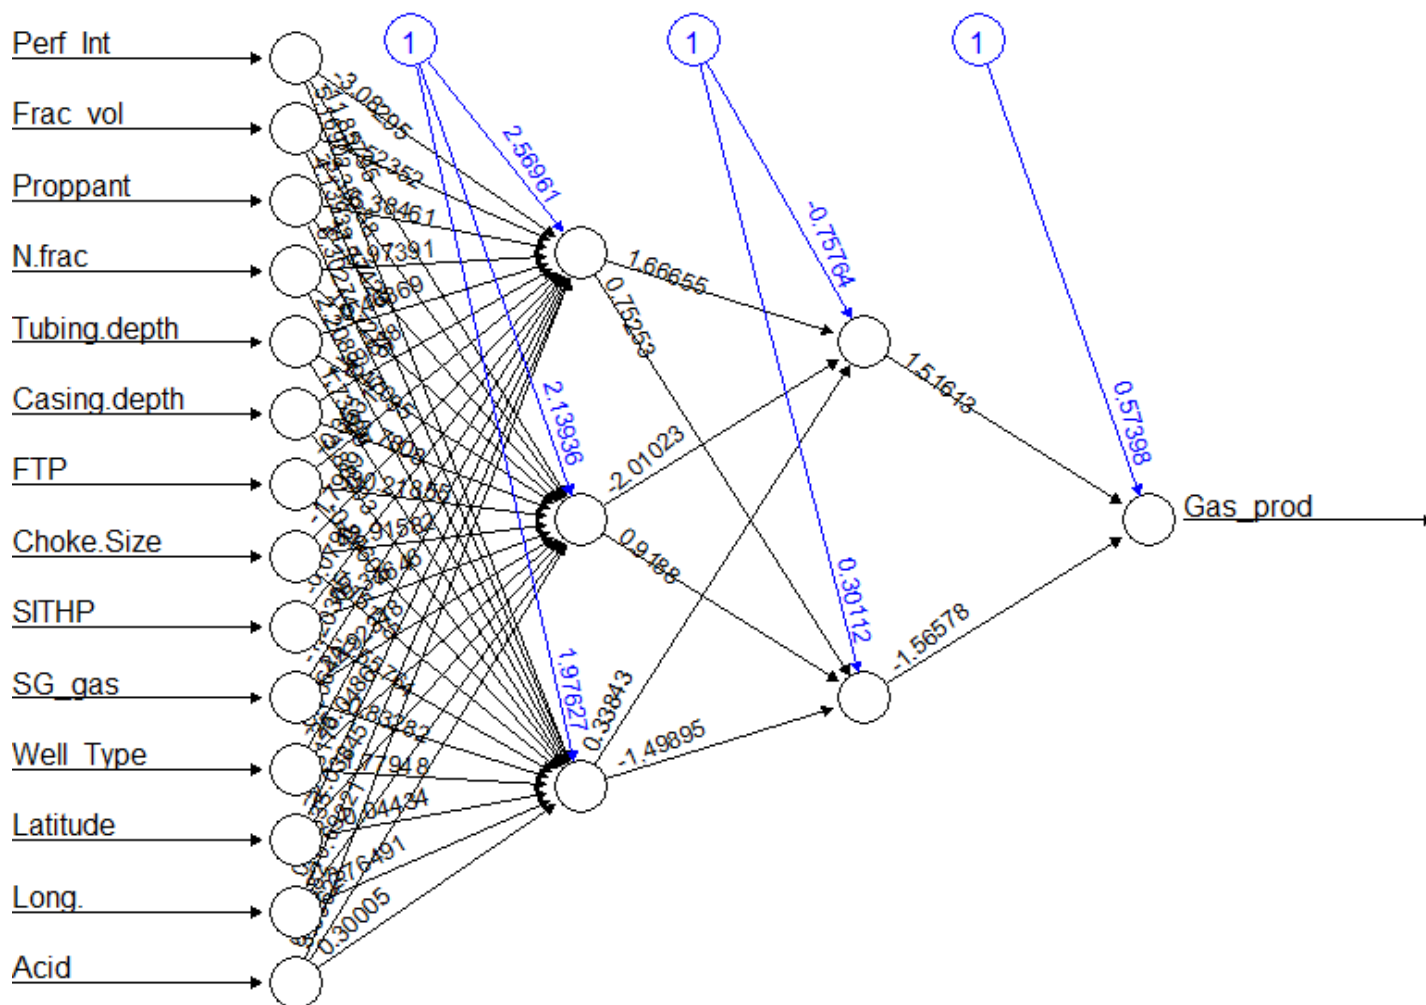
Hide

```
# Training and Test Data
trainset <- maxmindf[1:312, ]
testset <- maxmindf[313:446, ]
```

We fit the network and arbitrarily decide on the number of hidden neurons. Deciding on the number of hidden layers in a neural network is not an exact science. In fact, there are instances where accuracy will likely be higher without any hidden layers. Therefore, trial and error plays a significant role in this process. One possibility is to compare how the accuracy of the predictions change as we modify the number of hidden layers. If we use 3 and 2 hidden layers, we obtain:

Hide

```
# Neural Network Training
library(neuralnet)
nn <- neuralnet(Gas_prod ~ Perf_Int + Frac_vol + Proppant + N.frac +
  Tubing.depth + Casing.depth + FTP + Choke.Size + SITHP +
  SG_gas + Well_Type + Latitude + Long. + Acid, data=trainset,
  hidden=c(3,2), linear.output=TRUE, threshold=0.01)
nn$result.matrix
# Visual plot of the model
plot(nn)
```



The output from the neural network training is shown above. Our neural network has been created using the training data. We then compare this to the test data to gauge the accuracy of the neural network forecast.

[Hide](#)

```
#Test the resulting output
temp_test <- subset(testset, select = c("Gas_prod", "Perf_Int", "Frac_vol",
  "Proppant", "N.frac", "Tubing.depth",
  "Casing.depth", "FTP", "Choke.Size",
  "SITHP", "SG_gas", "Well_Type", "Latitude",
  "Long.", "Acid"))

head(temp_test)
```

	Gas_prod<dbl>	Perf_Int<dbl>	Frac_vol<dbl>	Proppant<dbl>	N.frac<dbl>	Tubing.depth<dbl>	Casing.depth<dbl>	
313	0.0115296075	0.14902047	0.006530593	0.08508562	0.0	0.4555120	0.2342327	C
314	0.0045662511	0.05414924	0.005322704	0.04983619	0.2	0.4542529	0.2495274	C
315	0.0041483181	0.14880035	0.005769774	0.05718191	0.0	0.3940963	0.2488400	C
316	0.0029264959	0.08056350	0.004923563	0.01567906	0.0	0.3567431	0.1031105	C
317	0.0009303225	0.11182038	0.007402423	0.07036233	0.3	0.4847510	0.3101908	C
318	0.0042217903	0.15100154	0.003430199	0.03379303	0.1	0.4956631	0.2758206	C

6 rows | 1-9 of 15 columns

Hide

```
nn.results <- compute(nn, temp_test)
```

The predicted results are compared to the actual results. The code to do this and a snippet of the result is shown below:

Hide

```
#Comparison of Predicted to Actual
results <- data.frame(actual = testset$Gas_prod, prediction = nn.results$net.result)
results
```

	actual<dbl>	prediction<dbl>
313	0.0115296075	0.0019106450
314	0.0045662511	0.0024498630
315	0.0041483181	0.0053135253
316	0.0029264959	0.0061688705
317	0.0009303225	0.0071988651
318	0.0042217903	0.0023350951
319	0.0052641792	0.0032005803
320	0.0065977434	0.0061051029
321	0.0033872906	0.0021572123
322	0.0051492189	-0.0008556349

1-10 of 134 rows

Previous123456...14Next

We then test the accuracy of the model

Hide

```
#Testing The Accuracy Of The Model
Gas_prod <- trainset$Gas_prod
predicted = results$prediction * abs(diff(range(Gas_prod))) + min(Gas_prod)
actual = results$actual * abs(diff(range(Gas_prod))) + min(Gas_prod)
comparison = data.frame(predicted, actual)

deviation = ((actual-predicted)/actual)
comparison = data.frame(predicted, actual, deviation)
```

In the above code, we are converting the data back to its original format. Note that we are also converting our data back into standard values given that they were previously scaled using the max-min normalization technique.

We compute the accuracy of network with (3,2) hidden layer. An accuracy of 61.8% on a mean absolute deviation basis (i.e. the average deviation between estimated and actual Gas production per month) was obtained.

Hide

```
accuracy = 1-abs(mean(deviation))
accuracy
```

```
[1] 0.6187621
```

You can see that we obtain a high accuracy of 61.8 % accuracy using a (3,2) hidden configuration. This is quite good, especially considering that our dependent variable is in the interval format. However, let's see if we can get it higher!

What happens if we now use a (5,2) hidden configuration in our neural network? Here is the generated output:

Hide

```
nn <- neuralnet(Gas_prod ~ Perf_Int + Frac_vol + Proppant + N.frac + Tubing.depth
                + Casing.depth + FTP + Choke.Size + SITHP + SG_gas + Well_Type
                + Latitude + Long. + Acid,data=trainset, hidden=c(5,2),
                linear.output=TRUE, threshold=0.01)
nn$result.matrix
```

```

[ ,1]
error 3.653272e-03
reached.threshold 8.863541e-03
steps 1.902000e+03
Intercept.to.1layhid1 2.412039e+00
Perf_Int.to.1layhid1 -3.022457e+00
Frac_vol.to.1layhid1 -2.706479e+00
Proppant.to.1layhid1 2.138009e+00
N.frac.to.1layhid1 6.900633e+00
Tubing.depth.to.1layhid1 3.020498e+00
Casing.depth.to.1layhid1 -1.606061e-01
FTP.to.1layhid1 -3.427725e-01
Choke.Size.to.1layhid1 2.963539e+00
SITHP.to.1layhid1 -7.742261e-02
SG_gas.to.1layhid1 -3.814676e-01
Well_Type.to.1layhid1 6.916603e-01
Latitude.to.1layhid1 -8.057662e-01
Long..to.1layhid1 -1.101212e+00
Acid.to.1layhid1 1.165234e+01
Intercept.to.1layhid2 -2.018957e+00
Perf_Int.to.1layhid2 4.331287e+00
Frac_vol.to.1layhid2 5.084424e-01
Proppant.to.1layhid2 -7.344221e+00
N.frac.to.1layhid2 -3.569281e+00
Tubing.depth.to.1layhid2 -5.241477e-01
Casing.depth.to.1layhid2 -1.099082e+00
FTP.to.1layhid2 -1.129401e+01
Choke.Size.to.1layhid2 3.903377e+00
SITHP.to.1layhid2 4.751302e+00
SG_gas.to.1layhid2 2.653925e+00
Well_Type.to.1layhid2 7.466284e+01
Latitude.to.1layhid2 -1.541364e+00
Long..to.1layhid2 9.906799e+00
Acid.to.1layhid2 1.626440e+00
Intercept.to.1layhid3 3.169506e+00
Perf_Int.to.1layhid3 -8.899430e+00
Frac_vol.to.1layhid3 -1.330377e+00
Proppant.to.1layhid3 8.086839e+01
N.frac.to.1layhid3 -6.402368e+00
Tubing.depth.to.1layhid3 -4.970220e-01
Casing.depth.to.1layhid3 -3.024804e+01
FTP.to.1layhid3 -1.423828e+01
Choke.Size.to.1layhid3 4.380193e+01
SITHP.to.1layhid3 1.046698e+01
SG_gas.to.1layhid3 9.484196e+00
Well_Type.to.1layhid3 5.987929e+01
Latitude.to.1layhid3 1.540507e-01
Long..to.1layhid3 -7.249328e+01
Acid.to.1layhid3 -8.130646e+00
Intercept.to.1layhid4 -3.390121e+00
Perf_Int.to.1layhid4 7.855518e+00
Frac_vol.to.1layhid4 7.786203e-01
Proppant.to.1layhid4 -2.010102e+01

```

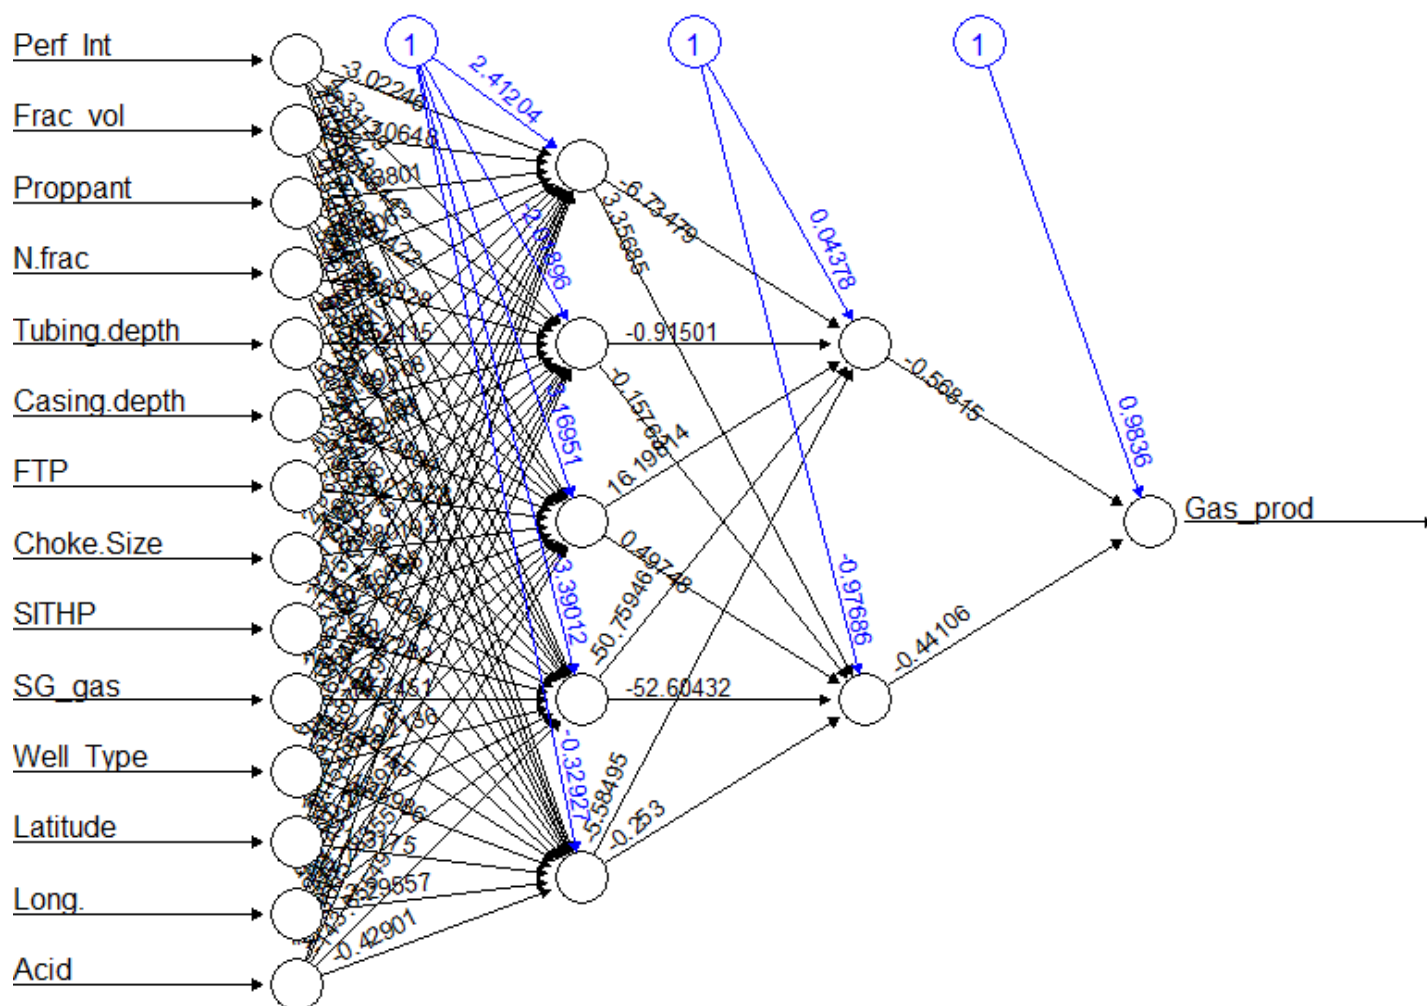
```

N.frac.to.1layhid4      6.809802e+00
Tubing.depth.to.1layhid4 4.636598e-01
Casing.depth.to.1layhid4 1.147943e+01
FTP.to.1layhid4         1.591844e+01
Choke.Size.to.1layhid4  -7.406063e+01
SITHP.to.1layhid4       -9.272816e+00
SG_gas.to.1layhid4      -1.074510e+01
Well_Type.to.1layhid4   -1.392136e+01
Latitude.to.1layhid4    2.439652e-01
Long..to.1layhid4       1.728355e+01
Acid.to.1layhid4        -1.436555e+02
Intercept.to.1layhid5   -3.292692e-01
Perf_Int.to.1layhid5    1.132030e+00
Frac_vol.to.1layhid5    -1.340956e+01
Proppant.to.1layhid5    3.907111e+00
N.frac.to.1layhid5      3.751044e+00
Tubing.depth.to.1layhid5 -2.037910e-01
Casing.depth.to.1layhid5 9.653307e-01
FTP.to.1layhid5         2.657104e+00
Choke.Size.to.1layhid5  -9.004901e-01
SITHP.to.1layhid5       -1.096745e+00
SG_gas.to.1layhid5      -2.615459e-01
Well_Type.to.1layhid5   -1.459861e+00
Latitude.to.1layhid5    1.317501e+00
Long..to.1layhid5       -2.295573e+00
Acid.to.1layhid5        -4.290123e-01
Intercept.to.2layhid1   4.378251e-02
1layhid1.to.2layhid1    -6.734787e+00
1layhid2.to.2layhid1    -9.150114e-01
1layhid3.to.2layhid1    1.619814e+01
1layhid4.to.2layhid1    -5.075946e+01
1layhid5.to.2layhid1    -5.584953e+00
Intercept.to.2layhid2   -9.768579e-01
1layhid1.to.2layhid2    3.356849e+00
1layhid2.to.2layhid2    -1.576312e-01
1layhid3.to.2layhid2    4.974827e-01
1layhid4.to.2layhid2    -5.260432e+01
1layhid5.to.2layhid2    -2.530036e-01
Intercept.to.Gas_prod   9.835952e-01
2layhid1.to.Gas_prod    -5.681484e-01
2layhid2.to.Gas_prod    -4.410568e-01

```

[Hide](#)

```
plot(nn)
```


[Hide](#)

```
#Test the resulting output
temp_test <- subset(testset, select = c("Gas_prod", "Perf_Int", "Frac_vol",
                                         "Proppant", "N.frac", "Tubing.depth",
                                         "Casing.depth", "FTP", "Choke.Size",
                                         "SITHP", "SG_gas", "Well_Type", "Latitude",
                                         "Long.", "Acid"))

head(temp_test)
```

	Gas_prod <dbl>	Perf_Int <dbl>	Frac_vol <dbl>	Proppant <dbl>	N.frac <dbl>	Tubing.depth <dbl>	Casing.depth <dbl>	
313	0.0115296075	0.14902047	0.006530593	0.08508562	0.0	0.4555120	0.2342327	C
314	0.0045662511	0.05414924	0.005322704	0.04983619	0.2	0.4542529	0.2495274	C
315	0.0041483181	0.14880035	0.005769774	0.05718191	0.0	0.3940963	0.2488400	C
316	0.0029264959	0.08056350	0.004923563	0.01567906	0.0	0.3567431	0.1031105	C
317	0.0009303225	0.11182038	0.007402423	0.07036233	0.3	0.4847510	0.3101908	C
318	0.0042217903	0.15100154	0.003430199	0.03379303	0.1	0.4956631	0.2758206	C

6 rows | 1-9 of 15 columns

Hide

```
nn.results <- compute(nn, temp_test)

results <- data.frame(actual = testset$Gas_prod, prediction = nn.results$net.result)
results
```

	actual <dbl>	prediction <dbl>
313	0.0115296075	0.008300917
314	0.0045662511	0.003234412
315	0.0041483181	0.004158014
316	0.0029264959	0.005043247
317	0.0009303225	0.005416683
318	0.0042217903	0.004526156
319	0.0052641792	0.006067041
320	0.0065977434	0.004446578
321	0.0033872906	0.003426161
322	0.0051492189	0.006288651

1-10 of 134 rows

Previous 1 2 3 4 5 6 ... 14 Next

The predicted results are compared to the actual results. We then test the accuracy of the model

Hide

```
#Testing The Accuracy Of The Model
predicted = results$prediction * abs(diff(range(Gas_prod))) + min(Gas_prod)
actual = results$actual * abs(diff(range(Gas_prod))) + min(Gas_prod)
comparison = data.frame(predicted,actual)
deviation = ((actual-predicted)/actual)
comparison = data.frame(predicted,actual,deviation)
```

And compute the accuracy of network with (5,2) hidden layer.

Hide

```
accuracy = 1-abs(mean(deviation))
accuracy
```

```
[1] 0.7784113
```

You can see that we obtain 77.84 % network accuracy using a (5,2) hidden configuration. We see that our accuracy rate has now increased to nearly 78 %, indicating that modifying the number of hidden nodes has enhanced our model! This shows that we can increase the accuracy of the network to predict to a higher value by increasing the number of hidden layers.

Model Interpretability

From the solution above, we can observe that neural networks resemble black boxes a lot: explaining their outcome is much more difficult than explaining the outcome of simpler model such as a linear model. Therefore, depending on the kind of application you need, you might want to take into account this factor too. Furthermore, as you have seen above, extra care will be needed to fit a neural network and small changes can lead to different results.

In addition, we showed that neural network is better at predicting "Gas_prod" than the linear regression model. A better mean squared error (MSE) value was obtained using neural network than linear model. Finally, it is possible for neural networks to be more accurate at prediction than regression; however, it will take trial and error of many different hidden layer configurations to get this better prediction.

Conclusion

This project was developed to illustrate the relationship between a dependent variable and several independent variables using **Artificial Neural Network**. We have been able to develop a way such that we were able to select the training, validation and testing datasets. We also investigated how varying number of hidden layers and neurons affect the artificial neural network results. The result using Artificial Neural Network was compared to result from using linear regression. We can conclude that the Artificial Neural Network is better at prediction than the linear regression.